







Francisco Muñoz-Martinez

francisco.munoz2@um.es

Universidad de Murcia (UM)

Attention: Tutorial is being recorded

Ager	nda
------	-----

Time (CET)	Time (ET)	Topic	Presenter
14:00 - 14:40	8:00 - 8:40	Flexible Accelerators	Tushar Krishna
14:40 – 15:10	8:40 – 9:10	Cycle accurate simulation and Overview of STONNE	José Luis Abellán
15:10 – 16:10	9:10 - 10:10	(Hands-on) STONNE Deep-Dive	Francisco Muñoz-Martínez
16:10 – 16:40	10:10 - 10:40	Coffee Break	
16:40 – 17:10	10:40 - 11:10	(Hands-on) STONNE Deep-Dive	Francisco Muñoz-Martínez
17:10 - 17:40	11:10 - 11:40	Dataflow exploration for Graph Neural Networks	Raveesh Garg
17:50 – 18:00	11:50 – 12:00	Roadmap for Future Development	Manuel Acacio

Tutorial Website https://stonne-simulator.github.io/ASPLOSTUT.html

includes agenda and STONNE/OMEGA installation instructions

Acknowledgements



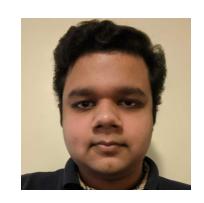
Tushar Krishna



José L. Abellán



Franciso Muñoz-Martínez



Raveesh Garg



Manuel E. Acacio

Outline

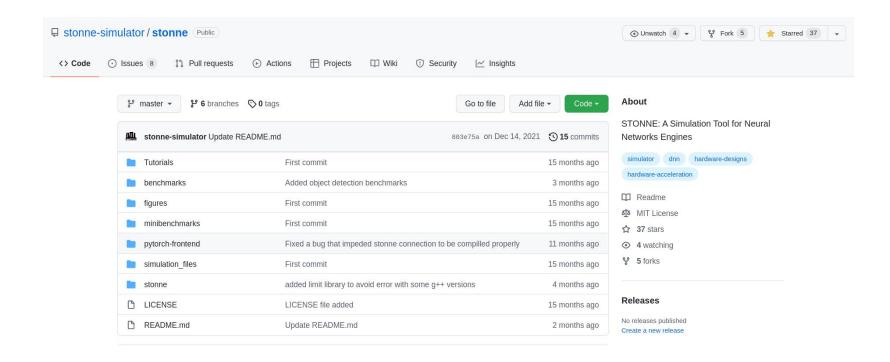
- Docker Image, Installation and overview of STONNE.
- Hands-on #1: Simulating a real DNN on Flexible DNN Accelerators.
- Hands-on #2: Using STONNE to evaluate a research use case Exploiting sparsity.
- Hands-on #3: Extending new operations in STONNE Adding the Conv1d operation.
- Research use cases.
- Conclusions.

Outline

- Docker Image, Installation and overview of STONNE.
- Hands-on #1: Simulating a real DNN on Flexible DNN Accelerators.
- Hands-on #2: Using STONNE to evaluate a research use case Exploiting sparsity.
- Hands-on #3: Extending new operations in STONNE Adding the Conv1d operation.
- Research use cases.
- Conclusions.

STONNE Repository

- STONNE available under the MIT license on Github:
 - https://github.com/stonne-simulator/stonne



STONNE Requirements

- STONNE dependencies:
 - Python 3.6 >= version < 3.9
 - C++ 14 or later.
 - Anaconda version 2021.05 or older.
 - .Linux OS (Tested on Ubuntu 18.05 and Manjaro 21.2.1).
 - Not tested in MAC OS.

STONNE Installation

- STONNE can be easily installed following a few steps:
 - STONNE User interface

```
(base) root@ae0d97074e84:/home/stonne omega/stonne# cd stonne/
(base) root@ae0d97074e84:/home/stonne omega/stonne/stonne# make
mkdir -p objs
                                -c src/DNNLayer.cpp -o objs/DNNLayer.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
q++ -03 -Iinclude/ -Iexternal/
                                -c src/Stats.cpp -o objs/Stats.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
                                -c src/OSMeshMN.cpp -o objs/OSMeshMN.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
                                -c src/CompilerMultiplierMesh.cpp -o objs/CompilerMultiplierMesh.o #-ltcmalloc
                                -c src/Accumulator.cpp -o objs/Accumulator.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
                                -c src/CompilerFEN.cpp -o objs/CompilerFEN.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
                                -c src/Fifo.cpp -o objs/Fifo.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
g++ -03 -Iinclude/ -Iexternal/
                                -c src/DataPackage.cpp -o objs/DataPackage.o #-ltcmalloc
                                -c src/TemporalRN.cpp -o objs/TemporalRN.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
                                -c src/CollectionBusLine.cpp -o objs/CollectionBusLine.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
q++ -03 -Iinclude/ -Iexternal/
                                -c src/MSwitch.cpp -o objs/MSwitch.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
                                -c src/MSNetwork.cpp -o objs/MSNetwork.o #-ltcmalloc
q++ -03 -Iinclude/ -Iexternal/
                                -c src/ASwitch.cpp -o objs/ASwitch.o #-ltcmalloc
                                -c src/Tile.cpp -o objs/Tile.o #-ltcmalloc
g++ -03 -Iinclude/ -Iexternal/
^CMakefile:23: recipe for target 'objs/Tile.o' failed
make: *** [objs/Tile.o] Interrupt
(base) root@ae0d97074e84:/home/stonne omega/stonne/stonne#
```

STONNE Installation

- STONNE can be easily installed following a few steps:
 - Pytorch-Frontend

```
(base) root@ae0d97074e84:/home/stonne omega/stonne# cd pytorch-frontend/
(base) root@ae0d97074e84:/home/stonne_omega/stonne/pytorch-frontend# python setup.py install
Building wheel torch-1.7.0a0
- Building version 1.7.0a0
cmake --build . --target install --config Release -- -j 8
  0%] Built target clog
  0%] Built target gtest
  0%] Built target defs.bzl
  0%] Built target pthreadpool
  0%] Built target benchmark
  1%] Built target libprotobuf-lite
  2%] Built target asmjit
  2%] Built target fmt
  3%] Built target c10
  3%] Built target foxi loader
  3%] Built target ATEN CPU FILES GEN TARGET
  4%] Built target dnnl common
  6%] Built target libprotobuf
  6%] Built target mkrename
  6%] Built target common
  6%] Built target mkdisp
  7%] Built target mkalias
 22%] Built target python copy files
 23%] Built target dnnl cpu
 23%] Built target arraymap
 23%] Built target mkrename_gnuabi
 23%] Built target mkmasked gnuabi
 23%] Built target torch global deps
 23%] Built target torch python stubs
```

STONNE Installation

- STONNE can be easily installed following a few steps:
 - Pytorch-stonne

```
(base) root@ae0d97074e84:/home/stonne omega/stonne/pytorch-frontend# cd stonne connection/
(base) root@ae0d97074e84:/home/stonne_omega/stonne/pytorch-frontend/stonne_connection# python setup.py install
['torch stonne.cpp', '../../stonne/stonne linker src/stonne linker.cpp', '../../stonne/src/Fifo.cpp', '../../st
nne/src/CollectionBusLine.cpp', '../../stonne/src/Connection.cpp', '../../stonne/src/FENetwork.cpp', '../../sto
ne/src/Accumulator.cpp', '../../stonne/src/MSwitch.cpp', '../../stonne/src/CompilerMultiplierMesh.cpp', '../..
tonne/src/DSNetworkTop.cpp', '../../stonne/src/DNNModel.cpp', '../../stonne/src/CompilerMSN.cpp', '../../stonne
src/OSMeshSDMemory.cpp', '../../stonne/src/DSwitch.cpp', '../../stonne/src/AccumulationBuffer.cpp', '../../ston
e/src/Confiq.cpp', '../../stonne/src/Tile.cpp', '../../stonne/src/Stats.cpp', '../../stonne/src/CollectionBus.
  ', '../../stonne/src/DNNLayer.cpp', '../../stonne/src/testbench.cpp', '../../stonne/src/STONNEModel.cpp', '..,
 /stonne/src/SDMemory.cpp', '../../stonne/src/MSNetwork.cpp', '../../stonne/src/LookupTable.cpp', '../../stonne
src/ASNetwork.cpp', '../../stonne/src/TemporalRN.cpp', '../../stonne/src/CompilerFEN.cpp', '../../stonne/src/M
tiplierOS.cpp', '../../stonne/src/DSNetwork.cpp', '../../stonne/src/DataPackage.cpp', '../../stonne/src/ASwitc
cpp', '../../stonne/src/OSMeshMN.cpp', '../../stonne/src/CompilerART.cpp', '../../stonne/src/SparseSDMemory.cpp
  '../../stonne/src/utility.cpp', '../../stonne/src/FEASwitch.cpp']
running install
running bdist egg
running egg info
writing torch stonne.egg-info/PKG-INFO
writing dependency links to torch stonne.egg-info/dependency links.txt
writing top-level names to torch stonne.egg-info/top level.txt
opt/conda/lib/python3.8/site-packages/torch/utils/cpp extension.py:340: UserWarning: Attempted to use ninja as/
the BuildExtension backend but we could not find ninja.. Falling back to using the slow distutils backend.
 warnings.warn(msg.format('we could not find ninja.'))
reading manifest file 'torch stonne.egg-info/SOURCES.txt'
writing manifest file 'torch stonne.egg-info/SOURCES.txt'
installing library code to build/bdist.linux-x86 64/egg
running install lib
```

• To avoid having to install the software, we have set up a docker image which can be directly downloaded and used:

sudo docker run -it franciscomunoz/stonne_omega_img/bin/bash

[paco@paco-pc stonne]\$ sudo docker run -it franciscomunoz/stonne_omega_img /bin/bash
(base) root@6e865a302d42:/home/stonne_omega/omega#

- Project organization:
 - /home/stonne_omega:
 - omega
 - stonne:
 - ASPLOS22
 - pytorch-frontend
 - stonne_connection
 - stonne:
 - stonne.elf
 - Include
 - src
 - external
 - stonne linker src

- Project organization:
 - /home/stonne_omega:
 - omega
 - stonne:
 - ASPLOS22
 - pytorch-frontend
 - stonne_connection
 - stonne:
 - stonne.elf
 - Include
 - src
 - external
 - stonne_linker_src

- Project organization:
 - /home/stonne_omega:
 - omega
 - stonne:
 - ASPLOS22
 - pytorch-frontend
 - stonne_connection
 - stonne:
 - stonne.elf
 - Include
 - src
 - external
 - stonne linker src

- Project organization:
 - /home/stonne_omega:
 - omega
 - stonne:
 - ASPLOS22
 - pytorch-frontend
 - stonne_connection
 - stonne:
 - stonne.elf
 - Include
 - src
 - external
 - stonne linker src

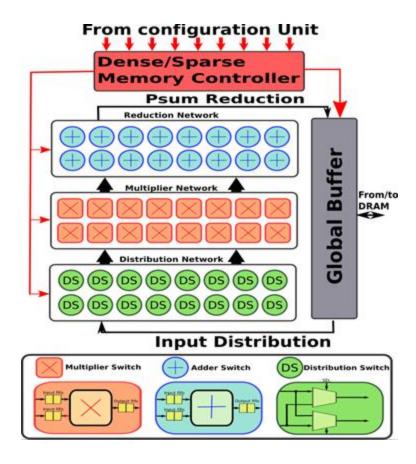
- Project organization:
 - /home/stonne_omega:
 - omega
 - stonne:
 - ASPLOS22
 - pytorch-frontend
 - stonne_connection
 - stonne:
 - stonne.elf
 - Include
 - src
 - external
 - stonne linker src

- Project organization:
 - /home/stonne_omega:
 - omega
 - stonne:
 - ASPLOS22
 - pytorch-frontend
 - stonne_connection
 - stonne:
 - stonne.elf
 - Include
 - src
 - external
 - stonne linker src


- Project organization:
 - /home/stonne_omega:
 - omega
 - stonne:
 - ASPLOS22
 - pytorch-frontend
 - stonne_connection
 - stonne:
 - stonne.elf
 - Include
 - src
 - external
 - stonne linker src

- /home/stonne_omega/stonne/stonne/include: Contains the headers
 - AccumulationBuffer.h
 - Accumulator.h
 - ReductionNetwork.h
 - ASNetwork.h
 - ASwitch.h
 - CollectionBus.h
 - CollectionBusLine.h
 - DistributionNetwork.h
 - DSNetworkTop.h
 - DSNetwork.h
 - DSwitch.h
 - MultiplierNetwork.h
 - MSNetwork.h
 - MultiplierOS.h
 - MSwitch.h
 - MemoryController.h
 - SDMemory.h
 - SparseSDMemory.h
 - OSMeshSDMemory.h
 - Connection.h
 - DataPackage.h
 - Fifo.h
 - STONNEModel.h

 STONNE organization is able to abstract the main components of the architecture



- Abstract classes allow to abtract the 4 main components.
 - AccumulationBuffer.h
 - Accumulator.h
 - ReductionNetwork.h
 - ASNetwork.h
 - ASwitch.h
 - CollectionBus.h
 - CollectionBusLine.h
 - DistributionNetwork.h
 - DSNetworkTop.h
 - DSNetwork.h
 - DSwitch.h
 - MultiplierNetwork.h
 - MSNetwork.h
 - MultiplierOS.h
 - MSwitch.h
 - MemoryController.h
 - SDMemory.h
 - SparseSDMemory.h
 - OSMeshSDMemory.h
 - Connection.h
 - DataPackage.h
 - Fifo.h
 - STONNEModel.h



- The real hardware comonents implement the abstract classes.
 - AccumulationBuffer.h
 - Accumulator.h
 - ReductionNetwork.h
 - ASNetwork.h
 - ASwitch.h
 - CollectionBus.h
 - CollectionBusLine.h
 - DistributionNetwork.h
 - DSNetworkTop.h 4
 - DSNetwork.h
 - DSwitch.h
 - MultiplierNetwork.h

 - MultiplierOS.h
 - MSwitch.h
 - MemoryController.h
 - SDMemory.h
 - SparseSDMemory.h
 - OSMeshSDMemory.h
 - Connection.h
 - DataPackage.h
 - Fifo.h
 - STONNEModel.h

Reduction Networks

Distribution Networks

MSNetwork.h

Multiplier Networks

Memory Controllers

Other Useful components

• Every component implements its **void cycle()** method which defines the behaviour of the component.

```
class MultiplierNetwork : public Unit{
public:
      By the default the implementation of the MS just receives a single element, calculate a single psum and/or send a single input ac
tivation to the neighbour. This way, the parameters
      input ports, output ports and forwarding ports will be set as the single data size. If this implementation change for future test
s, this can be change easily bu mofifying these three parameters.
   MultiplierNetwork(id t id, std::string name) : Unit(id, name){}
   //set connections from the distribution network to the multiplier network
   virtual void setInputConnections(std::map<int, Connection*> input connections) {assert(false);}
   //Set connections from the Multiplier Network to the Reduction Network
   virtual void setOutputConnections(std::map<int, Connection*> output connections) {assert(false);}
   virtual void cycle() {assert(false);}
   virtual void configureSignals(Tile* current tile, DNNLayer* dnn layer, unsigned int ms size, unsigned int n folding) {assert(false);
   virtual void configureSparseSignals(std::vector<SparseVN> sparseVNs, DNNLayer* dnn layer, unsigned int ms size) {assert(false);}
   virtual void resetSignals() {assert(false);}
   virtual void printConfiguration(std::ofstream& out, unsigned int indent) {assert(false);}
   virtual void printStats(std::ofstream &out, unsigned int indent) {assert(false);}
   virtual void printEnergy(std::ofstream& out, unsigned int indent) {assert(false);}
```

- /home/stonne_omega/stonne/stonne/src: Contains the main code
 - AccumulationBuffer.cpp
 - Accumulator.cpp
 - ASNetwork.cpp
 - ASwitch.cpp
 - CollectionBus.cpp
 - CollectionBusLine.cpp
 - DSNetworkTop.cpp
 - DSNetwork.cpp
 - DSwitch.cpp
 - MSNetwork.cpp
 - MultiplierOS.cpp
 - MSwitch.cpp
 - MemoryController.cpp
 - SDMemory.cpp
 - SparseSDMemory.cpp
 - OSMeshSDMemory.cpp
 - Connection.cpp
 - DataPackage.cpp
 - Fifo.cpp
 - STONNEModel.cpp

- STONNEModel.cpp drives the simulation based on the user configuration file
 - AccumulationBuffer.cpp
 - Accumulator.cpp
 - ASNetwork.cpp
 - ASwitch.cpp
 - CollectionBus.cpp
 - CollectionBusLine.cpp
 - DSNetworkTop.cpp
 - DSNetwork.cpp
 - DSwitch.cpp
 - MSNetwork.cpp
 - MultiplierOS.cpp
 - MSwitch.cpp
 - MemoryController.cpp
 - SDMemory.cpp
 - SparseSDMemory.cpp
 - OSMeshSDMemory.cpp
 - Connection.cpp
 - DataPackage.cpp
 - Fifo.cpp
 - STONNEModel.cpp



- STONNEModel.cpp drives the simulation based on the user configuration file
 - 1. The code reads the input file and selects the correct modules:

```
Stonne::Stonne(Config stonne cfg) {
   this->stonne_cfg=stonne_cfg;
   this->ms_size = stonne_cfg.m_MSNetworkCfg.ms_size;
   this->layer loaded=false:
   this->tile_loaded=false;
   this->outputASConnection = new Connection(stonne cfg.m SDMemoryCfg.port width);
   this->outputLTConnection = new Connection(stonne_cfg.m_LookUpTableCfg.port_width);
   switch(stonne_cfg.m_MSNetworkCfg.multiplier_network_type) {
       case LINEAR:
           this->msnet = new MSNetwork(2, "MSNetwork", stonne_cfg);
           break;
       case OS MESH:
           this->msnet = new OSMeshMN(2, "OSMesh", stonne cfg);
       default:
           assert(false);
   //switch(DistributionNetwork). It is possible to create instances of other Distributi
   this->dsnet = new DSNetworkTop(1, "DSNetworkTop", stonne_cfg);
   //Creating the ReduceNetwork according to the parameter specified by the user
   switch(stonne_cfg.m_ASNetworkCfg.reduce_network_type) {
   case ASNETWORK:
       this->asnet = new ASNetwork(3, "ASNetwork", stonne_cfg, outputASConnection);
   case FENETWORK:
       this->asnet = new FENetwork(3, "FENetwork", stonne_cfg, outputASConnection);
       break;
   case TEMPORALRN:
       this->asnet = new TemporalRN(3, "TemporalRN", stonne_cfg, outputASConnection);
       break;
   default:
       assert(false);
```

- STONNEModel.cpp drives the simulation based on the user configuration file
 - 2. After this, the class defines methods to load the layer, tile parameters, memory addresses and run the simulation. **This is used to interface the simulator**

```
void loadDNNLayer(Layer_t layer_type, std::string layer_name, unsigned int R, unsigned int S, unsigned int C, unsigned int K, unsigned int G, 
   N, unsigned int X, unsigned int Y, unsigned int strides, address_t input_address, address_t filter_address, address_t output_address, Dataflow dataflow);
General constructor
     //Load CONV Layer. At the end this calls to the general constructor with all the parameters
      void loadCONVLayer(std::string layer_name, unsigned int R, unsigned int S, unsigned int C, unsigned int K, unsigned int G, unsigned int N, unsigned int X
   unsigned int Y, unsigned int strides, address_t input_address, address_t filter_address, address_t output_address);
      //Load FC layer just with the appropriate parameters
      //N = batch size (i.e., number of rows in input matrix); S=number of inputs per batch (i.e., column size in input matrix and weight matrix); K=number of
   tputs neurons (i.e, number of rows weight matrix)
      void loadFCLayer(std::string layer_name, unsigned int N, unsigned int S, unsigned int K, address_t input_address, address_t filter_address, address_t outp
   t_address);
      //Load Sparse GEMM onto STONNE according to SIGMA parameter taxonomy.
      void loadGEMM(std::string layer_name, unsigned int N, unsigned int K, unsigned int M, address_t MK_matrix, address_t KN_matrix, metadata_address_t MK_meta
data, metadata_address_t KN_metadata, address_t output_matrix, metadata_address_t output_metadata, Dataflow dataflow);
     //Load Dense GEMM onto STONNE according to SIGMA parameter taxonomy and tiling according to T_N, T_K and T_M
     void loadDenseGEMM(std::string layer_name, unsigned int N, unsigned int K, unsigned int M, address_t MK_matrix, address_t KN_matrix, address_t output matri
    Dataflow dataflow):
      //Load sparse-dense GEMM onto STONNE
      void loadSparseDense(std::string layer_name, unsigned int N, unsigned int K, unsigned int M, address_t MK_matrix, address_t KN_matrix, metadata_address_t
MK metadata id, metadata address t MK metadata pointer, address t output matrix, unsigned int T N, unsigned int T K);
     //Load a Dense GEMM tile to run it using the loadDenseGEMM function
     void loadGEMMTile(unsigned int T N, unsigned int T K, unsigned int T M);
      void loadTile(unsigned int T_R, unsigned int T_S, unsigned int T_C, unsigned int T_K, unsigned int T_G, unsigned int T_N, unsigned int T_X_, unsigned int
   Y ); //Load general and CONV tile
      void loadFCTile(unsigned int T S, unsigned int T N, unsigned int T K); //VNSize = T S, NumVNs= T N*T K
       void run():
```

- STONNEModel.cpp drives the simulation based on the user configuration file
 - 3. The code runs the **cycle()** method of each component until the execution is completed.

```
bool execution finished=false;
while(!execution finished) {
    this->mem->cycle();
    this->collectionBus->cycle();
    this->asnet->cycle();
    this->lt->cycle();
      this->collectionBus->cycle(); //This order since these are connecti
    this->msnet->cycle();
    this->dsnet->cycle();
    execution finished = this->mem->isExecutionFinished();
    this->n cycles++;
if(this->stonne cfg.print stats enabled) { //If sats printing is enable
    this->printStats();
    this->printEnergy();
```

Outline


- Docker Image, Installation and overview of STONNE.
- Hands-on #1: Simulating a real DNN on Flexible DNN Accelerators.
- Hands-on #2: Using STONNE to evaluate a research use case Exploiting sparsity.
- Hands-on #3: Extending new operations in STONNE Adding the Conv1d operation.
- Research use cases.
- Conclusions.

Objectives


- 1. Comprehensively understand how do flexible DNN architectures work.
- 2. Understand how to run simulations in STONNE using its STONNE User interface.
- Understand how to run a real DNN in STONNE using the framework Pytorch.

DNN benchmark


- We will go through a real-life example to motivate our use case.
- Application task: Image digit recognition
- Benchmark: A Convolutional Neural Network (CNN).
- Dataset: MNIST (60000 28x28-size images)



CNN model




CNN model



Simulated architecture

• We will simulate our benchmark in a 64-MS MAERI architecture.



	Number of Multipliers		
1	64	64	64

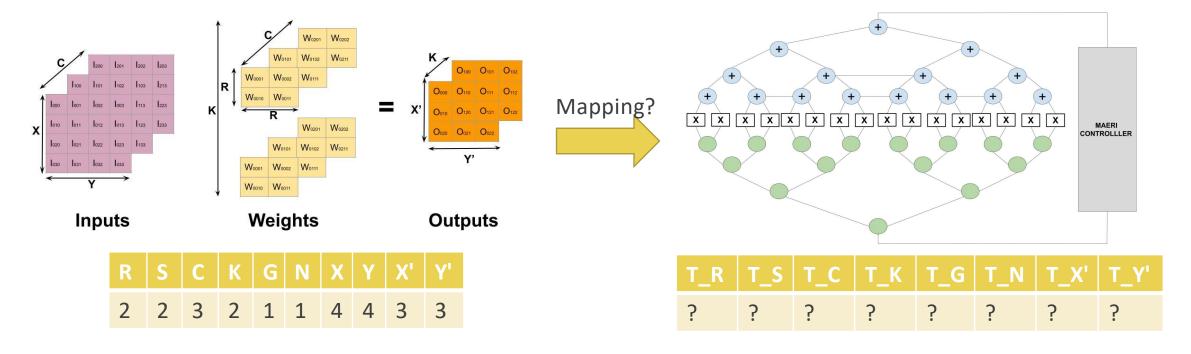
Mapping the CNN onto the architectures

• Given our CNN and our two architectures, how do we map the computation?

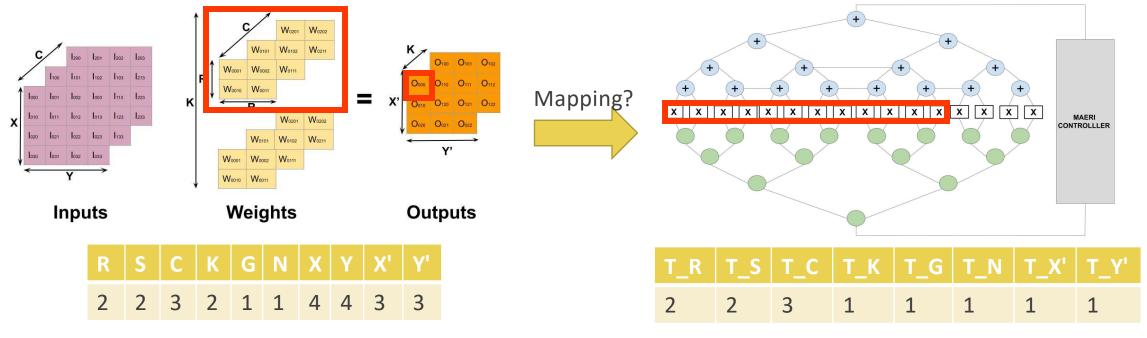
Mapping the CNN onto the architectures

- Given our CNN and our two architectures, how do we map the computation?
 - Answer: Tiling the layer.

- Given our CNN and our two architectures, how do we map the computation?
- Example: How do we map this layer?



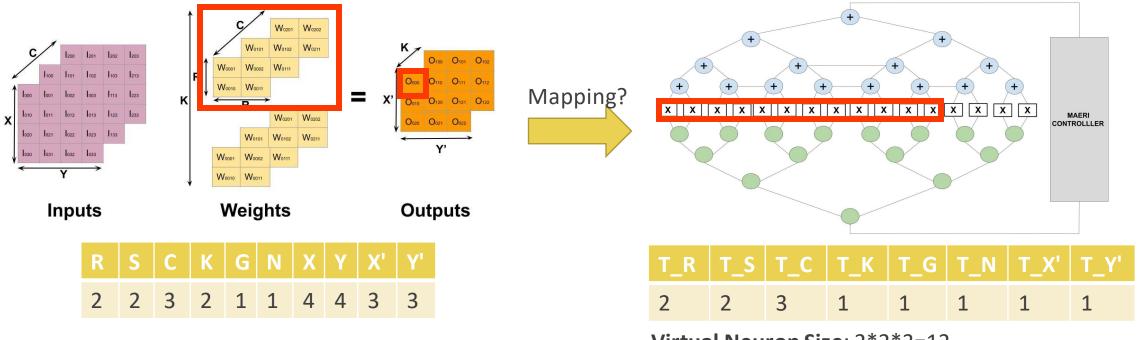
- Given our CNN and our two architectures, how do we map the computation?
- Example: How do we map this layer?



Virtual Neuron Size: ?

Number of Virtual Neurons:?

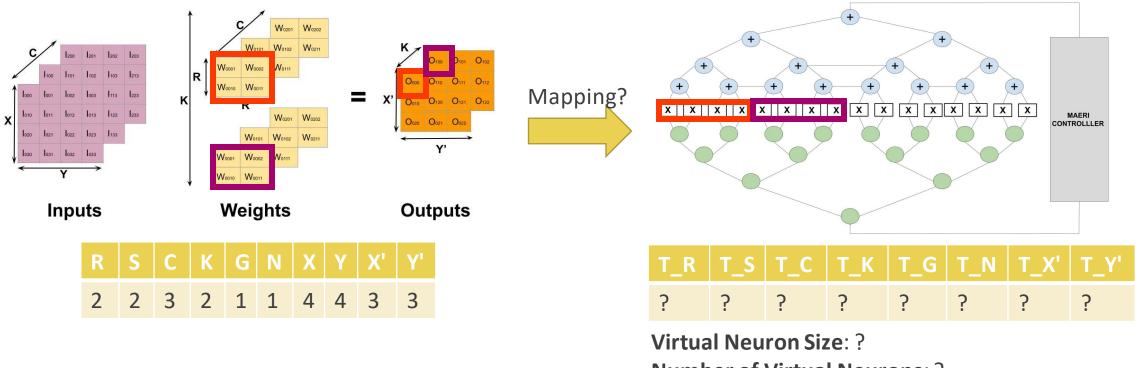
- Given our CNN and our two architectures, how do we map the computation?
- Example: How do we map this layer?



Virtual Neuron Size: 2*2*3=12

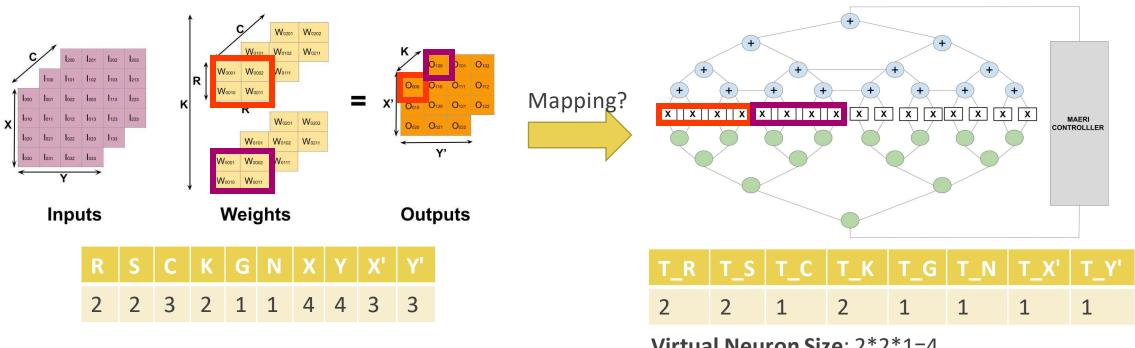
Number of Virtual Neurons: 1*1*1*1=1

- Given our CNN and our two architectures, how do we map the computation?
- Example: How do we map this layer?



Number of Virtual Neurons:?

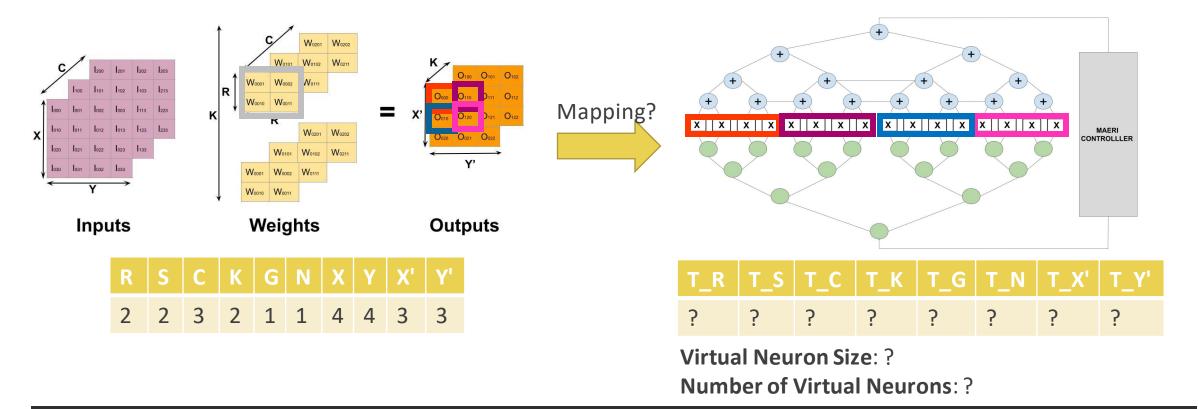
- Given our CNN and our two architectures, how do we map the computation?
- Example: How do we map this layer?



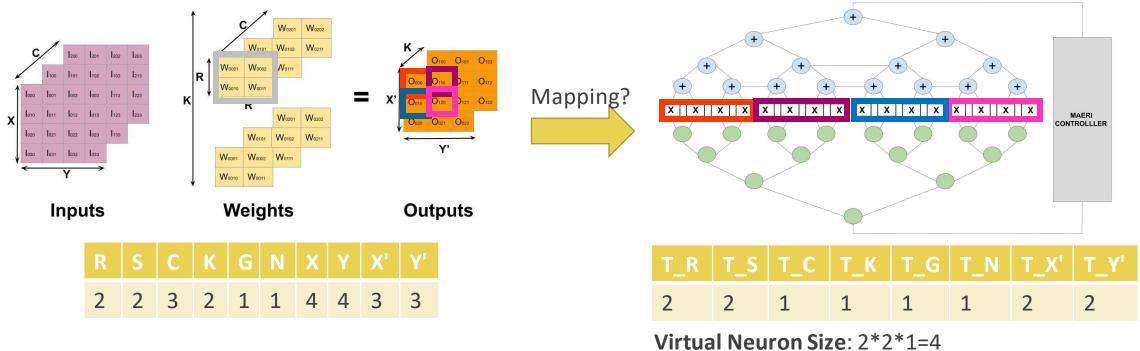
Virtual Neuron Size: 2*2*1=4

Number of Virtual Neurons: 2*1*1*1*1=2

- Given our CNN and our two architectures, how do we map the computation?
- Example: How do we map this layer?

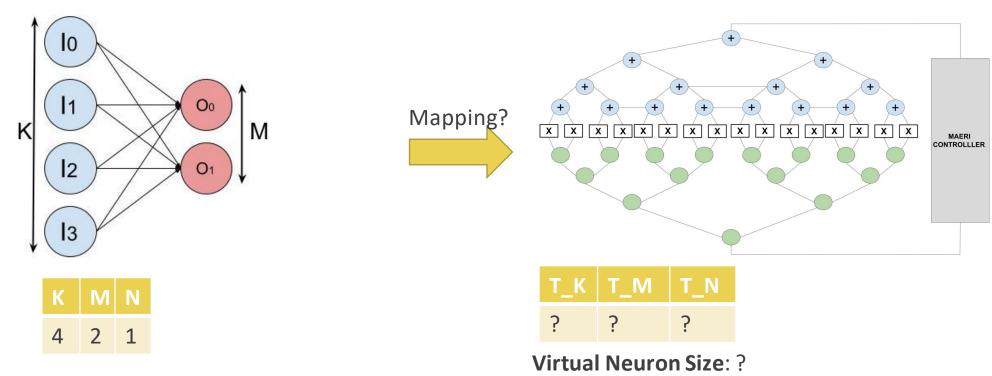


- Given our CNN and our two architectures, how do we map the computation?
- Example: How do we map this layer?



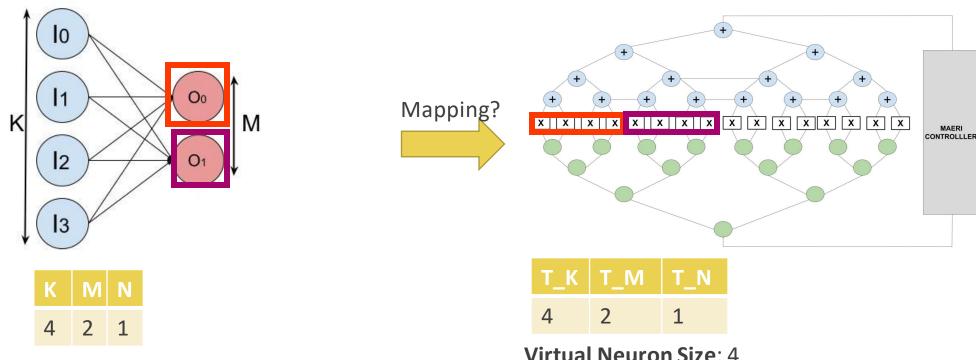
Number of Virtual Neurons: 1*1*1*2*2=4

- Given our CNN and our two architectures, how do we map the computation?
- Example: How do we map this layer?



Number of Virtual Neurons:?

- Given our CNN and our two architectures, how do we map the computation?
- Example: How do we map this layer?



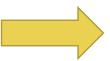
Virtual Neuron Size: 4

Number of Virtual Neurons: 2*1=2

• Given our CNN and our architecture, how do we map the computation?

Lay	yer	Operation	R	S	С	K	G	N	X	Υ				
0		CONV	5	5	1	16	1	1	32	32	Simulated	Number of	DN	RN
1		MaxPool	2	2	16	16	1	1	28	28		Multipliers	BW	BW
2		CONV	5	5	16	32	1	1	18	18	Simulated	64	64	64
3		MaxPool	2	2	32	32	1	1	14	14				
Lay	yer	Operation	K	M	N									
4		Dense	1568	10	1		-	-	-	-	Simulated			

Layer	Operation	R	S	С	K	G	N	X	Υ	
0	CONV	5	5	1	16	1	1	32	32	Simulated
1	MaxPool	2	2	16	16	1	1	28	28	
2	CONV	5	5	16	32	1	1	18	18	Simulated
3	MaxPool	2	2	32	32	1	1	14	14	
Layer	Operation	K	M	N						
4	Dense	1568	10	1		-	-	-	-	Simulated



Number of	DN	RN
Multipliers	BW	BW
64	64	64

Layer	Operation	T_R	T_S	T_C	T_K	T_G	T_N	T_X'	T_Y'	VN Size	# VNs	MSs used
0	CONV	?	?	?	?	?	?	?	?	?	?	?
2	CONV	?	?	?	?	?	?	?	?	?	?	?
Layer	Operation	T_K	T_M	T_N								

Dense

Layer	Operation	R	S	С	K	G	N	X	Υ					
0	CONV	5	5	1	16	1	1	32	32	Simu	lated			Num
1	MaxPool	2	2	16	16	1	1	28	28					Multi
2	CONV	5	5	16	32	1	1	18	18	Simu	lated			64
3	MaxPool	2	2	32	32	1	1	14	14					
Layer	Operation	K	M	N										
4	Dense	1568	10	1		-	-	-	-	Simu	llated			
Layer	Operation	T_R	T S	T	СТ	K	T_(5 T	N	T_X'	T_Y'	VN	# VNs	MSs u

Number of	DN	RN
Multipliers	BW	BW
64	64	64

Layer	Operation	T_R	T_S	T_C	T_K	T_G	T_N	T_X'	T_Y'	VN Size	# VNs	MSs used
0	CONV	5	5	1	2	1	1	1	1	25	2	50
2	CONV	?	?	?	?	?	?	?	?	?	?	?
Layer	Operation	T_K	T_M	T_N								
4	Dense	?	?	?	?	?	?	?	?	?	?	?

Layer	Operation	R	S	С	K	G	N	X	Υ					
0	CONV	5	5	1	16	1	1	32	32	Simu	ılated			Num
1	MaxPool	2	2	16	16	1	1	28	28					Mult
2	CONV	5	5	16	32	1	1	18	18	Simu	lated			64
3	MaxPool	2	2	32	32	1	1	14	14					
Layer	Operation	K	M	N										
4	Dense	1568	10	1		-	-	-	-	Simu	ılated			
Layer	Operation	T_R	T_S	T_0	C T	_K	T_0	G T_	N	T_X'	T_Y'	VN	# VNs	MSs

64	64	64
MSs used		

nber of

tipliers

50

64

25

16

1

16

T_M

0

Layer

CONV

CONV

Dense

Operation

BW

BW

Layer	Operation	R	S	С	K	G	N	X	Υ					
0	CONV	5	5	1	16	1	1	32	32	Simu	ılated			Num
1	MaxPool	2	2	16	16	1	1	28	28					Mult
2	CONV	5	5	16	32	1	1	18	18	Simu	lated			64
3	MaxPool	2	2	32	32	1	1	14	14					
Layer	Operation	K	M	N										
4	Dense	1568	10	1		-	-	-	-	Simu	ılated			
Layer	Operation	T_R	T_S	T_	C T	_K	T_0	G T_	_N	T_X'	T_Y'	VN	# VNs	MSs

4	Delise	1308	10	1			_					
Layer	Operation	T_R	T_S	T_C	T_K	T_G	T_N	T_X'	T_Y'	VN Size	# VNs	MSs used
0	CONV	5	5	1	2	1	1	1	1	25	2	50
2	CONV	1	1	16	4	1	1	1	1	16	4	64
Layer	Operation	T_K	T_M	T_N								
4	Dense	32	2	1	-	-	-	-	-	32	2	64

nber of

tipliers

BW

64


BW

64

- Before running the real benchmark we will use the STONNE User interface to simulate our layers.
- This will run the layers with the correct dimensions but using random data.

```
[PC@User $] ./stonne -CONV -R=3 -S=3 -C=1 -G=1 -K=1 -N=1 -X=4 -Y=4 -T_R=3 -T_S=3 -T_C=1 -T_G=1 -T_K=1 -T_N=1 -T_X_=1 -T_Y_=1 -num_ms=64 -dn_bw=64 -rn_bw=64 [PC@User $] Running CONV layer in STONNE Simulator...
[PC@User $] Output files generated correctly.
```

 Once the simulation has finished we can get the energy and area numbers:



- Script to calculate the energy and area numbers:
 - /home/stonne_omega:
 - omega
 - stonne:
 - ASPLOS22
 - pytorch-frontend
 - stonne_connection
 - stonne:
 - stonne.elf
 - Include
 - src
 - external
 - stonne_linker_src
 - energy_tables
 - calculate_energy.py
 - energy_model.txt


Syntax: ./calculate_energy [-v] -table_file=<Energy&Area
 table> -counter_file=<counter stats> -out_file=<output file

- Our image digit recognition benchmark is located in /home/stonne_omega/stonne/ASPLOS22/handson-1
 - train_dnn.py: Used to train our CNN and get the trained weights.
 - evaluate_dnn.py: Used to evaluate the accuracy of our CNN.
 - dnn.py: Defines our CNN.
 - deployment.py: This is the final file which is used to predict real images.

- Steps towards CPU execution:
 - 1. Run train_dnn.py script to obtain the weights.
 - 2. Run evaluate_dnn.py to measure the accuracy of the network.
 - 3. Run deployment.py

- Steps towards CPU execution:
 - 1. Run train_dnn.py script to obtain the weights.
 - 2. Run evaluate_dnn.py to measure the accuracy of the network.
 - 3. Run deployment.py

- Once we have tested the inference procedure on the CPU, let's run it in our simulated accelerator.
- Steps towards simulation:
 - 1. Set up the hardware configuration file.
 - 2. Set up the tile configuration files
 - 3. Change the CONV and Linear operations in order to simulate them.



Outline


- Docker Image, Installation and overview of STONNE.
- Hands-on #1: Simulating a real DNN on Flexible DNN Accelerators.
- Hands-on #2: Using STONNE to evaluate a research use case Exploiting sparsity.
- Hands-on #3: Extending new operations in STONNE Adding the Conv1d operation.
- Research use cases.
- Conclusions.


Objectives

- 1. Understand why running the simulation with the real numbers is of paramount importance.
- 2. Understand how to modify the hardware components in the STONNE simulator.
- Understand why STONNE can be useful to research.

Presence of zeros


• DNNs present a large number of zeros due to both prunning (weights) and the ReLU activation function (activations).






Gate operations


- Some accelerators leverage this feature to save energy by skiping the computation involving the zeros:
 - Any multiplication involving a zero is zero.
 - Output = i*w + p



Compression

- Other accelerators compress the zeros (e.g., using a CSR format or bitmap) and avoid useless transfers from memory.
- e.g., SIGMA:






Research use case

- In this use case we will evaluate the impact of the gate operation technique when running our benchmark.
- Sparse executions will be explored in OMEGA.
- Gate operations (i.e., avoid computing zeros) are not yet implemented.

Research use case

The idea is to create another multiplier network with multipliers implementing this feature.



Research use case


- ¿How many operations do we compute?
- ¿How can we add a new stat tracking the number of operations that are skiped?
- ¿What is the percentage of operations that are skiped?
- ¿What implications does it have on both performance and energy?

Outline

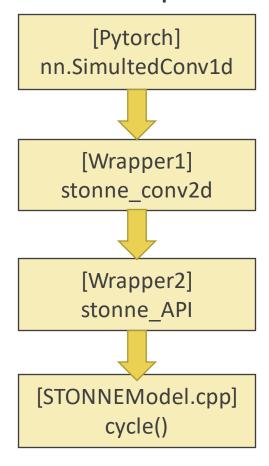
- Docker Image, Installation and overview of STONNE.
- Hands-on #1: Simulating a real DNN on Flexible DNN Accelerators.
- Hands-on #2: Using STONNE to evaluate a research use case Exploiting sparsity.
- Hands-on #3: Extending new operations in STONNE –
 Adding the Conv1d operation.
- Research use cases.
- Conclusions.

- At this point, we have seen how the operations nn.Linear and nn.Conv2d can be simulated in STONNE.
- What if we want to integrate a new operation and how does the Pytorch-STONNE connection work?

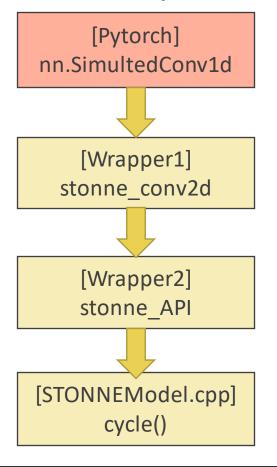
• The nn.conv1d operation is not integrated with STONNE.



• During this exercise, we will see how to connect the conv1d operation to the convolution operation implemented in STONNE.

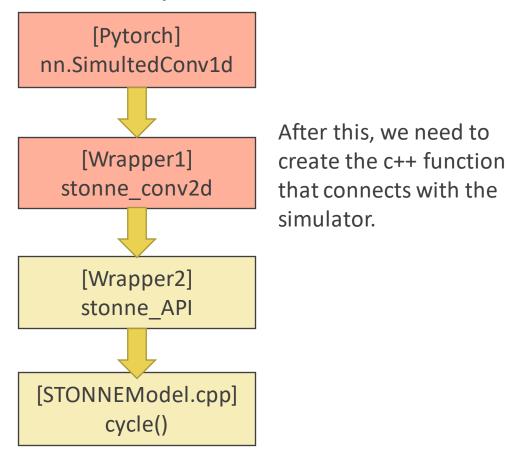


• During this exercise, we will see how to connect the conv1d operation to the convolution operation implemented in STONNE.

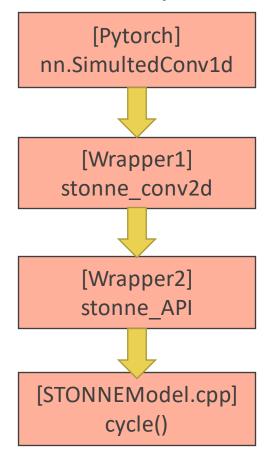


First we will create the operation within the **nn** package in Pytorch.

• During this exercise, we will see how to connect the conv1d operation to the convolution operation implemented in STONNE.

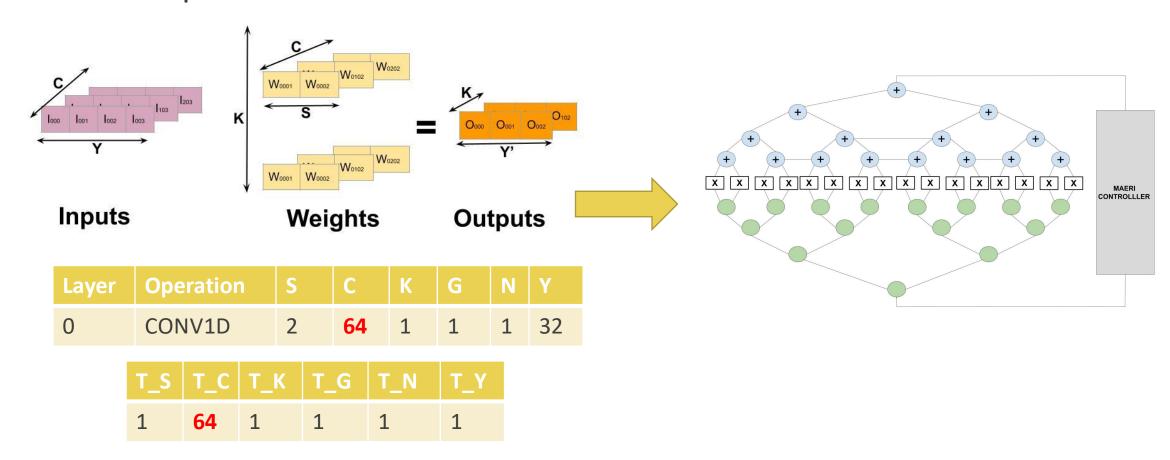


• During this exercise, we will see how to connect the conv1d operation to the convolution operation implemented in STONNE.



The STONNE back-end remains intact as the simulator already supports the convolution operation.

• After implementing the operation we will use it to run the next CONV1D operation on our 64-MS architecture:

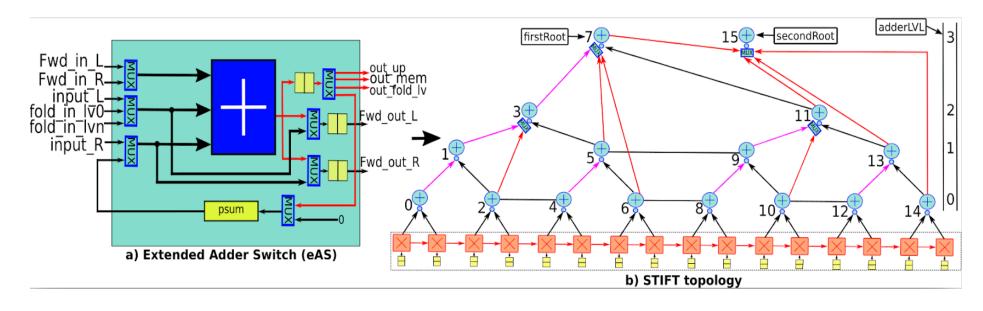


Outline

- Docker Image, Installation and overview of STONNE.
- Hands-on #1: Simulating a real DNN on Flexible DNN Accelerators.
- Hands-on #2: Using STONNE to evaluate a research use case Exploiting sparsity.
- Hands-on #3: Extending new operations in STONNE Adding the Conv1d operation.
- Research use cases.
- Conclusions.

Use case #1: STIFT

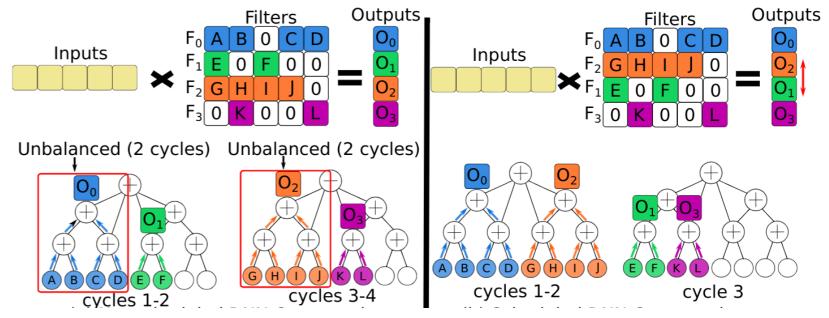
• We can use STONNE to create and evaluate new architectural modules for ML accelerators.



STIFT [Francisco Munoz-Martinez et. al., NOCS21]

Use case #2: Sparse filters scheduling

• We can use STONNE to evaluate new compilation techniques such as sparse scheduling.



STONNE [Francisco Munoz-Martinez et. al., IISWC21]

Outline

- Docker Image, Installation and overview of STONNE.
- Hands-on #1: Simulating a real DNN on Flexible DNN Accelerators.
- Hands-on #2: Using STONNE to evaluate a research use case Exploiting sparsity.
- Hands-on #3: Extending new operations in STONNE Adding the Conv1d operation.
- Research use cases.
- Conclusions.

Conclusions

- STONNE is a cycle-accurate simulator for ML accelerators.
- During this hands-on we have seen:
 - How to use the STONNE user interface to run synthetic operations.
 - How to use the Pytorch interface to run end-to-end models with STONNE.
 - How to extend the architecture of STONNE.
 - How to connect new operations and frameworks to the STONNE simulator.



Find STONNE on http://github.com/stonne simulator/stonne













Francisco Muñoz-Martínez

francisco.munoz2@um.es

Universidad de Murcia (UM)